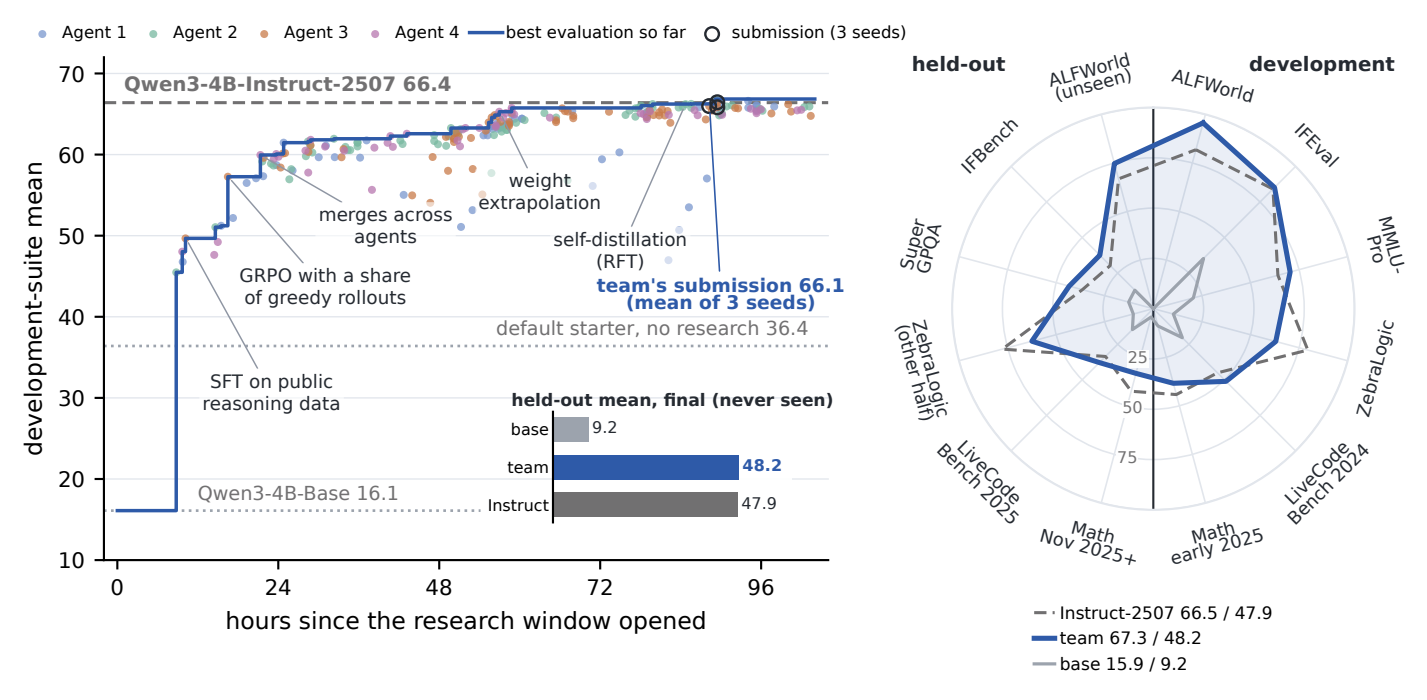
Agent 1
Agent 2
Agent 3
Agent 4
best evaluation so far
submission (3 seeds)
Qwen3-4B-Instruct-2507 66.4
development-suite mean
hours since the research window opened
weight extrapolation
merges across agents
self-distillation (RFT)
GRPO with a share of greedy rollouts
team's submission 66.1 (mean of 3 seeds)
default starter, no research 36.4
SFT on public reasoning data
Qwen3-4B-Base 16.1
held-out mean, final (never seen)
base 9.2
team 48.2
Instruct 47.9
held-out
development
ALFWorld (unseen)
ALFWorld
IFBench
IFEval
Super GPQA
MMLU-Pro
ZebraLogic (other half)
ZebraLogic
LiveCode Bench 2025
LiveCode Bench 2024
Math Nov 2025+
Math early 2025
Instruct-2507 66.5 / 47.9
team 67.3 / 48.2
base 15.9 / 9.2